

- Step 1
- Specify user details
- Step 2
- Set permissions
- Step 3
- Review and create
- Step 4
- Retrieve password

Specify user details

User details

User name

Sonai

The user name can have up to 64 characters. Valid characters: A-Z, a-z, 0-9, and + , . , @ _ - (hyphen)

☒ Provide user access to the AWS Management Console - optional

In addition to console access, users with SignInLocalDevelopmentAccess permissions can use the same console credentials for programmatic access without the need for access keys.

Console password

☐ Autogenerated password

You can view the password after you create the user.

☒ Custom password

Enter a custom password for the user.

- Must be at least 8 characters long
- Must include at least three of the following mix of character types: uppercase letters (A-Z), lowercase letters (a-z), numbers (0-9), and symbols ! @ # \$ % ^ & * () _ + - (hyphen) = [] { } ' "

☐ Show password

☒ Users must create a new password at next sign-in - Recommended

Users automatically get the IAMUserChangePassword policy to allow them to change their own password.

Set permissions

Add user to an existing group or create a new one. Using groups is a best-practice way to manage user's permissions by job functions. [Learn more](#)

Permissions options

☐ Add user to group

Add user to an existing group, or create a new group. We recommend using groups to manage user permissions by job function.

☐ Copy permissions

Copy all group memberships, attached managed policies, and inline policies from an existing user.

☒ Attach policies directly

Attach a managed policy directly to a user. As a best practice, we recommend attaching policies to a group instead. Then, add the user to the appropriate group.

Permissions policies (20/1445)

Choose one or more policies to attach to your new user.



Create policy

Filter by Type

Search

All types

<

1

2

3

4

5

6

7

...

73

>

☒	Policy name	Type	Attached entities
☒	 AccessAnalyzerServiceRolePolicy	AWS managed	0
☒	 AccountManagementFromVercel	AWS managed	0
☒	 AdministratorAccess	AWS managed - job function	0
☒	 AdministratorAccess-Amplify	AWS managed	0
☒	 AdministratorAccess-AWSElasticBeanstalk	AWS managed	0
☒	 AIOpsAssistantIncidentReportPolicy	AWS managed	0
☒	 AIOpsAssistantPolicy	AWS managed	0
☒	 AIOpsConsoleAdminPolicy	AWS managed	0
☒	 AIOpsOperatorAccess	AWS managed	0
☒	 AIOpsReadOnlyAccess	AWS managed	0

Review and create

Review your choices. After you create the user, you can view and download the autogenerated password, if enabled.

User details

User name

Sonai

Console password type

Custom password

Require password reset


Yes

Permissions summary

< 1 2 3 >

Name	Type	Used as
AccessAnalyzerServiceRolePolicy	AWS managed	Permissions policy
AccountManagementFromVercel	AWS managed	Permissions policy
AdministratorAccess	AWS managed - job function	Permissions policy
AdministratorAccess-Amplify	AWS managed	Permissions policy
AdministratorAccess-AWSElasticBeanstalk	AWS managed	Permissions policy
AIOpsAssistantIncidentReportPolicy	AWS managed	Permissions policy
AIOpsAssistantPolicy	AWS managed	Permissions policy
AIOpsConsoleAdminPolicy	AWS managed	Permissions policy
AIOpsOperatorAccess	AWS managed	Permissions policy
AIOpsReadOnlyAccess	AWS managed	Permissions policy

System Architecture:



Screenshots of Static Environment:

EC2 > Instances > Launch an instance

Launch an instance

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags

Name

Static_Instance

Add additional tags

Application and OS Images (Amazon Machine Image)

An AMI contains the operating system, application server, and applications for your instance. If you don't see a suitable AMI below, use the search field or choose Browse more AMIs.

Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Amazon Linux

aws

macOS

Mac

Ubuntu

ubuntu

Windows

Microsoft

Red Hat

Red Hat

SUSE Linux

SUSE

Debian

debian

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

Summary

Number of instances

1

Software Image (AMI)

Amazon Linux 2023 AMI 2023.10...[read more](#)

ami-0c1fe732b5494dc14

Virtual server type (instance type)

t2.micro

Firewall (security group)

New security group

Storage (volumes)

1 volume(s) - 8 GiB

Cancel

Launch instance

Preview code

Instance type

Info | Get advice

Instance type

t2.micro

Family: t2 1 vCPU 1 GiB Memory Current generation: true

On-Demand Linux base pricing: 0.0116 USD per Hour

On-Demand Windows base pricing: 0.0162 USD per Hour

On-Demand Ubuntu Pro base pricing: 0.0134 USD per Hour

On-Demand SUSE base pricing: 0.0116 USD per Hour On-Demand RHEL base pricing: 0.026 USD per Hour

All generations

Compare instance types

Additional costs apply for AMIs with pre-installed software

Key pair (login)

Info

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

Sonai

Create new key pair

Network settings

Info

VPC - required

Info

vpc-02492443d8bcb7b0e (default)

Create new VPC

Subnet

Info

No preference

Create new subnet

Availability Zone

Info

No preference

Enable additional zones

Auto-assign public IP

Info

Enable

Firewall (security groups)

Info

Create security group

Select existing security group

Security group name - required

Static_Env

This security group will be added to all network interfaces. The name can't be edited after the security group is created. Max length is 255 characters. Valid characters: a-z, A-Z, 0-9, spaces, and _-:/@.#[]+=&()*~*

Description - required

Info

launch-wizard-6 created 2026-02-11T16:37:43.651Z

Inbound Security Group Rules

Security group rule 1 (TCP; 22; 0.0.0.0/0)

Remove

Type

Info

ssh

Protocol

Info

TCP

Port range

Info

22

Source type

Info

Anywhere

Source

Info

0.0.0.0/0

Description - optional

Info

e.g. SSH for admin desktop

Security group rule 2 (TCP; 80; 0.0.0.0/0)

Remove

Type

Info

HTTP

Protocol

Info

TCP

Port range

Info

80

Source type

Info

Anywhere

Source

Info

0.0.0.0/0



Description - optional

Info

e.g. SSH for admin desktop

Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only.

Add security group rule



Instance summary for i-01f8ee4e4cb2f8668 (Static_Instance) [Info](#)

Updated 3 minutes ago

Instance ID
 i-01f8ee4e4cb2f8668

IPv6 address
–

Hostname type
IP name: ip-172-31-23-138.ec2.internal

Answer private resource DNS name
IPv4 (A)

Auto-assigned IP address
 3.91.60.155 [Public IP]

IAM role
–

IMDSv2
Required

Operator

Public IPv4 address
 3.91.60.155 | [open address](#)

Instance state
 Running

Private IP DNS name (IPv4 only)
 ip-172-31-23-138.ec2.internal

Instance type
t2.micro

VPC ID
 vpc-02492443d8bcb7b0e

Subnet ID
 subnet-0b1657fa09eff134b

Instance ARN
 arn:aws:ec2:us-east-1:001187921517:instance/i-01f8ee4e4cb2f8668

[Connect](#) [Instance state](#) [Actions](#)

Private IPv4 addresses
 172.31.23.138

Public DNS
 ec2-3-91-60-155.compute-1.amazonaws.com | [open address](#)

Elastic IP addresses
–

AWS Compute Optimizer finding
 [Opt-in to AWS Compute Optimizer for recommendations.](#) | [Learn more](#)

Auto Scaling Group name
–

Managed
false

[EC2](#) > [Volumes](#) > Create volume

Create volume [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

Volume settings

Volume type | [Info](#)

General Purpose SSD (gp3)

Size (GiB) | [Info](#)

8

Min: 1 GiB, Max: 65536 GiB.

IOPS | [Info](#)

3000

Min: 3000 IOPS, Max: 80000 IOPS.

Throughput (MiB/s) | [Info](#)

125

Min: 125 MiB, Max: 2000 MiB. Baseline: 125 MiB/s.

Availability Zone | [Info](#)

use1-az4 (us-east-1c)

Snapshot ID - optional | [Info](#)

Don't create volume from a snapshot



Encryption | [Info](#)

Use Amazon EBS encryption as an encryption solution for your EBS resources associated with your EC2 instances.

☐ Encrypt this volume

[Alt+S]

Swagata Sinha (0013-8792-1517)

Swagata Sinha

Instances (1/1) Info

Find Instance by attribute or tag (case-sensitive)

All states

Static_Instance

i-01f8ee4e4cb2f8668

Running

t2.micro

2/2 checks passed

View alarms +

us-east-1c

ec2-3-91-60-155.com...

Instance diagnostics

Instance settings

Networking

Security

Image and templates

Storage

Monitor and troubleshoot

Attach volume

Detach volume

[Alt+S]

Swagata Sinha

EC2 > Instances > i-01f8ee4e4cb2f8668 > Attach volume

Attach volume info

You can create and configure EBS volumes in your account and then attach them to EC2 instances in your VPC.

Instance ID

i-01f8ee4e4cb2f8668 (Static_Instance)

Availability Zone

us-east-1c

Currently attached volumes (1)

Volume

vol-0d80b816eb56a0147

Create new volume

Device name

/dev/sdb

Recommended device names for Linux: /dev/xvda for root volume, /dev/sd(f-p) for data volumes.

Cancel Attach volume

[Alt+S]

Swagata Sinha

Instances (1/1) Info

Find Instance by attribute or tag (case-sensitive)

All states

Static_Instance

i-01f8ee4e4cb2f8668

Running

t2.micro

2/2 checks passed

View alarms +

us-east-1c

ec2-3-91-60-155.comp...

3.91.60.155

-

Detailed monitoring

After you enable detailed monitoring for an instance, monitoring data is available in 1-minute periods. [Learn more](#)

Instance ID

i-01f8ee4e4cb2f8668 (Static_Instance)

Detailed monitoring

☒ Enable

After you enable detailed monitoring, the Amazon EC2 console displays monitoring graphs with a 1-minute period for the instance. [Additional charges apply](#)

Cancel Confirm

Details

Status and alarms

Monitoring

Security

Instance summary info

Instance ID

i-01f8ee4e4cb2f8668

Public IPv4 addresses

72.31.23.138

i-01f8ee4e4cb2f8668 (Static_Instance)

Details

Status and alarms

Monitoring

Security

Networking

Storage

Tags

☒ Include metrics in the CWAgent namespace [Learn more](#)

☐ Alarm recommendations [Investigate with AI - new](#) 1h

CPU utilization (%)

Percent

2.71%

1.35%

0%

18:00 18:15 18:30 18:45

Network in (bytes)

Bytes

4.224k

2.112k

0

18:00 18:15 18:30 18:45

[Alt+S]

Swagata Sinha

EC2 > Instances > i-01f8ee4e4cb2f8668 > Connect to instance

Connect info

Connect to an instance using the browser-based client.

Instance ID

i-01f8ee4e4cb2f8668 (Static_Instance)

VPC ID

vpc-02492443d8bcb7b0e

Security groups

sg-0ea7005259cd9c29e (Static_Env)

IAM role

-

EC2 Instance Connect

SSM Session Manager

SSH client

EC2 serial console

Instance ID

i-01f8ee4e4cb2f8668 (Static_Instance)

Connection type

☒ Connect using a Public IP

☐ Connect using a Private IP

Connect using a public IPv4 or IPv6 address

Connect using a private IP address and a VPC endpoint

Public IPv4 address

3.91.60.155

IPv6 address

-

Username


ec2-user

Note: In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel Connect

☒ Include metrics in the CWAgent namespace ⓘ
[Learn more](#) ↗

Investigate with AI - new 1h



Screenshots of Dynamic ASG Environment:

Launch an instance

Amazon EC2 allows you to create virtual machines, or instances, that run on the AWS Cloud. Quickly get started by following the simple steps below.

Name and tags

Name

Dynamic_Instance

Add additional tags

Application and OS Images (Amazon Machine Image)

An AMI contains the operating system, application server, and applications for your instance. If you don't see a suitable AMI below, use the search field or choose [Browse more AMIs](#).

Search our full catalog including 1000s of application and OS images

Recents

Quick Start

Amazon Linux

macOS

Ubuntu

Windows

Red Hat

SUSE Linux

Debian

Browse more AMIs

Including AMIs from AWS, Marketplace and the Community

Amazon Machine Image (AMI)

Amazon Linux 2023 kernel-6.1 AMI
ami-0c1fe732b5494dc14 (64-bit (x86), uefi-preferred) / ami-02b86da1e539b4dd0 (64-bit (ARM), uefi)
Virtualization: hvm ENA enabled: true Root device type: ebs

Description

Amazon Linux 2023 (kernel-6.1) is a modern, general purpose Linux-based OS that comes with 5 years of long term support. It is optimized for AWS and designed to provide a secure, stable and high-performance execution environment to develop and run your cloud applications.

Summary

Number of instances1

Software Image (AMI)
Amazon Linux 2023 AMI 2023.10...[read more](#)
ami-0c1fe732b5494dc14

Virtual server type (instance type)
t2.micro

Firewall (security group)
Static_Erv

Storage (volumes)
1 volume(s) - 8 GiB

Cancel

Launch instance

Preview code

Instance type

Instance type

t2.micro

Family: t2 1 vCPU 1 GiB Memory Current generation: true On-Demand Linux base pricing: 0.0116 USD per Hour
On-Demand Windows base pricing: 0.0162 USD per Hour On-Demand Ubuntu Pro base pricing: 0.0134 USD per Hour
On-Demand SUSE base pricing: 0.0116 USD per Hour On-Demand RHEL base pricing: 0.026 USD per Hour

All generations

[Compare instance types](#)

Additional costs apply for AMIs with pre-installed software

Key pair (login)

You can use a key pair to securely connect to your instance. Ensure that you have access to the selected key pair before you launch the instance.

Key pair name - required

Sonali

[Create new key pair](#)

Network settings

Network

vpc-02492443d8bcb7b0e

Subnet

No preference (Default subnet in any availability zone)

Auto-assign public IP

Enable

Firewall (security groups)

A security group is a set of firewall rules that control the traffic for your instance. Add rules to allow specific traffic to reach your instance.

Create security group

Select existing security group

Common security groups

Select security groups

Static_Erv sg-0ea7005259cd9c29e

VPC: vpc-02492443d8bcb7b0e

Security groups that you add or remove here will be added to or removed from all your network interfaces.

[Compare security group rules](#)

Configure storage

1x 8 GiB gp3 Root volume, 3000 IOPS, Not encrypted

[Add new volume](#)

Click refresh to view backup information

The tags that you assign determine whether the instance will be backed up by any Data Lifecycle Manager policies.

Summary

Number of instances

1

Software Image (AMI)

Amazon Linux 2023 AMI 2023.10...[read more](#)
ami-0c1fe732b5494dc14

Virtual server type (instance type)

t2.micro

Firewall (security group)

Static_Erv

Storage (volumes)

1 volume(s) - 8 GiB

Cancel

Launch instance

[Preview code](#)

Success

Successfully initiated launch of instance (i-01a166f1641ae4756)

Launch log

Initializing requests Succeeded

Launch initiation Succeeded

Instance summary for i-01a166f1641ae4756 (Dynamic_Instance) Info

Updated 5 minutes ago

Instance ID

i-01a166f1641ae4756

Public IPv4 address

3.80.87.24 | [open address ↗](#)

Private IPv4 addresses

172.31.21.186

IPv6 address

—

Instance state

Running

Public DNS

ec2-3-80-87-24.compute-1.amazonaws.com | [open address ↗](#)

Hostname type

IP name: ip-172-31-21-186.ec2.internal

Answer private resource DNS name

IPv4 (A)

Auto-assigned IP address

3.80.87.24 [Public IP]

IAM role

—

IMDSv2

Required

Operator

—

Detailed monitoring

✕

After you enable detailed monitoring for an instance, monitoring data is available in 1-minute periods. [Learn more ↗](#)

Instance ID

i-01a166f1641ae4756 (Dynamic_Instance)

Detailed monitoring

☒ Enable

After you enable detailed monitoring, the Amazon EC2 console displays monitoring graphs with a 1-minute period for the instance. [Additional charges apply ↗](#)


Cancel

Confirm

Details | Status and alarms | **Monitoring** | Security | Networking | Storage | Tags

- ☒ Include metrics in the CWAgent namespace Info
[Learn more ↗](#)
- ☐ Alarm recommendations Info
- Investigate with AI - new

1h



[Instances](#) > [i-01a166f1641ae4756](#) > [Connect to instance](#)

Connect Info

Connect to an instance using the browser-based client.

Instance ID

i-01a166f1641ae4756 (Dynamic_Instance)

VPC ID

vpc-02492443d8bcb7b0e

Security groups

sg-0ea7005259cd9c29e (Static_Env)

IAM role

—

EC2 Instance Connect

SSM Session Manager

SSH client

EC2 serial console

Instance ID

i-01a166f1641ae4756 (Dynamic_Instance)

Connection type

☒ Connect using a Public IP

Connect using a public IPv4 or IPv6 address

☐ Connect using a Private IP

Connect using a private IP address and a VPC endpoint

☒ Public IPv4 address

3.80.87.24

☐ IPv6 address

—

Username

Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

ec2-user

✕

Note: In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

Connect

aws

Search

[Alt+S]

United States (N. Virginia)

Swagata Sinha

```
Amazon Linux 2023
https://aws.amazon.com/linux/amazon-linux-2023

[ec2-user@ip-172-31-21-186 ~]$ pwd
/home/ec2-user
[ec2-user@ip-172-31-21-186 ~]$ sudo yum install stress -y

Amazon Linux 2023 Kernel Livepatch repository
Last metadata expiration check: 0:00:01 ago on Wed Feb 11 19:42:18 2026.
Dependencies resolved.

Package      Architecture Version      Repository Size
Installing:
stress       x86_64      1.0.7-2.amzn2023.0.1 amazonlinux 34 k

Transaction Summary
Install 1 Package

Total download size: 34 k
Installed size: 68 k
Downloading Packages:
stress-1.0.7-2.amzn2023.0.1.x86_64.rpm
1.0 MB/s | 34 kB | 00:00
Running transaction check
Transaction check succeeded.
Running transaction test
Transaction test succeeded.
Running transaction
Preparing :
Installing : stress-1.0.7-2.amzn2023.0.1.x86_64
1/1
1/1
```

i-01a166f1641ae4756 (Dynamic_Instance)

Create volume [Info](#)

Create an Amazon EBS volume to attach to any EC2 instance in the same Availability Zone.

Volume settings

Volume type [Info](#)

General Purpose SSD (gp3) ▼

Size (GiB) [Info](#)

8

Min: 1 GiB, Max: 65536 GiB.

IOPS [Info](#)

3000

Min: 3000 IOPS, Max: 80000 IOPS.

Throughput (MiB/s) [Info](#)

125

Min: 125 MiB, Max: 2000 MiB. Baseline: 125 MiB/s.

Availability Zone [Info](#)

use1-az4 (us-east-1c) ▼

Snapshot ID - optional [Info](#)

Don't create volume from a snapshot ▼ ↺

Encryption [Info](#)

Use Amazon EBS encryption as an encryption solution for your EBS resources associated with your EC2 instances.

☐ Encrypt this volume

[Alt+S] Ask Amazon Q

United States (N. Virginia) Swagata Sinha

Instances (1/1) [Info](#)

Last updated less than a minute ago

Connect Instance state ▼ Actions ▲ Launch instances ▼

Find Instance by attribute or tag (case-sensitive) All states ▼

Instance state: running X Clear filters

☒	Name ↗	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
☒	Dynamic_Insta...	i-01a166f1641ae4756	Running	t2.micro	2/2 checks passed	View alarms +	us-east-1c	ec2-3-80-87-24.com

Instance diagnostics

Instance settings

Networking

Security

Image and templates

Storage

Monitor and troubleshoot

Search [Alt+S] Ask Amazon Q

United States (N. Virginia) Swagata Sinha

Instances > i-01a166f1641ae4756 > Attach volume

Swagata Sinha

Attach volume [Info](#)

You can create and configure EBS volumes in your account and then attach them to EC2 instances in your VPC.

Instance ID
i-01a166f1641ae4756 (Dynamic_Instance)

Availability Zone
us-east-1c

▶ Currently attached volumes (1)

Volume [Info](#)

vol-0435006e19d0b23d8

Create new volume ↗

Device name [Info](#)

/dev/sdb

Recommended device names for Linux: /dev/xvda for root volume, /dev/sd[f-p] for data volumes.

Cancel Attach volume



i-01a166f1641ae4756 (Dynamic_Instance)

[Learn more](#)

☒ Alarm recommendations

[Investigate with AI - new](#)

1h



[Alt+S] Ask Amazon Q

United States (N. Virginia) Swagata Sinha

Elastic-Storage-ASG-Template

Elastic-Storage-ASG-Template (lt-08c8c1bb9146f7956)

Actions Delete template

Launch template details

Launch template ID
lt-08c8c1bb9146f7956

Launch template name
Elastic-Storage-ASG-Template

Default version
1

Owner
arnawsiam:001187921517:root

Details

Versions

Template tags

Launch template version details

Actions Delete template version

Version
1 (Default)

Description
-

Date created
2026-02-11T20:03:47.000Z

Created by
arnawsiam:001187921517:root

Instance details

Storage

Resource tags

Network interfaces

Advanced details

AMI ID
ami-0b7ef51a683137062

Instance type
t2.micro

Availability Zone
-

Availability Zone Id
-

Key pair name
Sonal

Security groups
-

Security group IDs
sg-0ea7005259cd9c29e

[Alt+S] Ask Amazon Q

United States (N. Virginia) Swagata Sinha

Scaling group

Choose launch template

Step 2

Choose instance launch options

Step 3 - optional

Integrate with other services

Step 4 - optional

Configure group size and scaling

Step 5 - optional

Add notifications

Step 6 - optional

Add tags

Step 7

Review

Choose launch template

Specify a launch template that contains settings common to all EC2 instances that are launched by this Auto Scaling group.

Name

Auto Scaling group name
Enter a name to identify the group.
Dynamic-CloudStorage-ASG
Must be unique to this account in the current Region and no more than 255 characters.

Launch template

For accounts created after May 31, 2023, the EC2 console only supports creating Auto Scaling groups with launch templates. Creating Auto Scaling groups with launch configurations is not recommended but still available via the CLI and API until December 31, 2023.

Launch template
Choose a launch template that contains the instance-level settings, such as the Amazon Machine Image (AMI), instance type, key pair, and security groups.
Elastic-Storage-ASG-Template
Create a launch template

Version
Default (1)
Create a launch template version

Description
-

Launch template
Elastic-Storage-ASG-Template
lt-08c8c1bb9146f7956

Instance type
t2.micro

AMI ID
ami-0b7ef51a683137062

Security groups
-

Request Spot Instances
No

Key pair name
Sonal

Security group IDs
sg-0ea7005259cd9c29e

Additional details

Storage (volumes)
-

Date created
Thu Feb 12 2026 01:33:47 GMT+0530 (India Standard Time)

Choose instance launch options

Choose the VPC network environment that your instances are launched into, and customize the instance types and purchase options.

Instance type requirements

Override launch template

You can keep the same instance attributes or instance type from your launch template, or you can choose to override the launch template by specifying different instance attributes or manually adding instance types.

Launch template
Elastic-Storage-ASG-Template
lt-08c8c1bb9146f7956

Version
Default

Description

Instance type
t2.micro

Network

For most applications, you can use multiple Availability Zones and let EC2 Auto Scaling balance your instances across the zones. The default VPC and default subnets are suitable for getting started quickly.

VPC

Choose the VPC that defines the virtual network for your Auto Scaling group.

vpc-02492443d8bcb7b0e
172.31.0.0/16 Default

Create a VPC

Availability Zones and subnets

Define which Availability Zones and subnets your Auto Scaling group can use in the chosen VPC.

Select Availability Zones and subnets

use1-az4 (us-east-1c) | subnet-0b1657fa09eff134b
172.31.16.0/20 Default

use1-az1 (us-east-1a) | subnet-0f6dd550b930586bd
172.31.0.0/20 Default

use1-az2 (us-east-1b) | subnet-0866de87910fe34ae
172.31.80.0/20 Default

use1-az6 (us-east-1d) | subnet-0cbb951ae534cf4ec
172.31.32.0/20 Default

Create a subnet

Availability Zone distribution - new

Auto Scaling automatically balances instances across Availability Zones. If launch failures occur in a zone, select a strategy.

Balanced best effort

If launches fail in one Availability Zone, Auto Scaling will attempt to launch in another healthy Availability Zone.

Balanced only

If launches fail in one Availability Zone, Auto Scaling will continue to attempt to launch in the unhealthy Availability Zone to preserve balanced distribution.

Cancel

Skip to review

Previous

Next

Integrate with other services - optional [Info](#)

Use a load balancer to distribute network traffic across multiple servers. Enable service-to-service communications with VPC Lattice. Shift resources away from impaired Availability Zones with zonal shift. You can also customize health check replacements and monitoring.

Load balancing [Info](#)

Use the options below to attach your Auto Scaling group to an existing load balancer, or to a new load balancer that you define.

Select Load balancing options

- ☒ No load balancer
Traffic to your Auto Scaling group will not be fronted by a load balancer.
- ☐ Attach to an existing load balancer
Choose from your existing load balancers.
- ☐ Attach to a new load balancer
Quickly create a basic load balancer to attach to your Auto Scaling group.

VPC Lattice integration options [Info](#)

To improve networking capabilities and scalability, integrate your Auto Scaling group with VPC Lattice. VPC Lattice facilitates communications between AWS services and helps you connect and manage your applications across compute services in AWS.

Select VPC Lattice service to attach

- ☒ No VPC Lattice service
VPC Lattice will not manage your Auto Scaling group's network access and connectivity with other services.
- ☐ Attach to VPC Lattice service
Incoming requests associated with specified VPC Lattice target groups will be routed to your Auto Scaling group.

[Create new VPC Lattice service](#) [↗](#)

Application Recovery Controller (ARC) zonal shift - new [Info](#)

During an Availability Zone impairment, target instance launches towards other healthy Availability Zones.

- ☐ Enable zonal shift
New instance launches will be retargeted towards healthy Availability Zones until the zonal shift is canceled.

Application Recovery Controller (ARC) zonal shift - new [Info](#)

During an Availability Zone impairment, target instance launches towards other healthy Availability Zones.

- ☐ Enable zonal shift
New instance launches will be retargeted towards healthy Availability Zones until the zonal shift is canceled.

Health checks

Health checks increase availability by replacing unhealthy instances. When you use multiple health checks, all are evaluated, and if at least one fails, instance replacement occurs.

EC2 health checks

[Always enabled](#)

Additional health check types - optional | [Info](#)

- ☐ Turn on Elastic Load Balancing health checks
Elastic Load Balancing monitors whether instances are available to handle requests. When it reports an unhealthy instance, EC2 Auto Scaling can replace it on its next periodic check.
- ☐ Turn on VPC Lattice health checks
VPC Lattice can monitor whether instances are available to handle requests. If it considers a target as failed a health check, EC2 Auto Scaling replaces it after its next periodic check.
- ☒ Turn on Amazon EBS health checks
EBS monitors whether an instance's root volume or attached volume stalls. When it reports an unhealthy volume, EC2 Auto Scaling can replace the instance on its next periodic health check.

[ⓘ](#) EBS volumes that are attached to an instance have a DeleteOnTermination attribute. When instances are terminated, this attribute either detaches the volumes (keeping data intact), or deletes the volumes. Detached volumes are not reattached when a new instance starts. To avoid data loss, take frequent snapshots of any critical EBS volumes. You can automate snapshot creation and deletion with [Amazon Data Lifecycle Manager](#) [↗](#)

[✕](#)

Health check grace period | [Info](#)

This time period delays the first health check until your instances finish initializing. It doesn't prevent an instance from terminating when placed into a non-running state.

seconds

[Cancel](#) [Skip to review](#) [Previous](#) [Next](#)

Configure group size and scaling - optional [Info](#)

Define your group's desired capacity and scaling limits. You can optionally add automatic scaling to adjust the size of your group.

Group size [Info](#)

Set the initial size of the Auto Scaling group. After creating the group, you can change its size to meet demand, either manually or by using automatic scaling.

Desired capacity type

Choose the unit of measurement for the desired capacity value. vCPUs and Memory(GiB) are only supported for mixed instances groups configured with a set of instance attributes.

Desired capacity

Specify your group size.

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

Equal or less than desired capacity

Max desired capacity

Equal or greater than desired capacity

Automatic scaling - optional

Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

- ☐ No scaling policies
Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.
- ☒ Target tracking scaling policy
Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling [Info](#)

You can resize your Auto Scaling group manually or automatically to meet changes in demand.

Scaling limits

Set limits on how much your desired capacity can be increased or decreased.

Min desired capacity

Equal or less than desired capacity

Max desired capacity

Equal or greater than desired capacity

Automatic scaling - *optional*

Choose whether to use a target tracking policy [Info](#)

You can set up other metric-based scaling policies and scheduled scaling after creating your Auto Scaling group.

☐ No scaling policies

Your Auto Scaling group will remain at its initial size and will not dynamically resize to meet demand.

☒ Target tracking scaling policy

Choose a CloudWatch metric and target value and let the scaling policy adjust the desired capacity in proportion to the metric's value.

Scaling policy name

Metric type [Info](#)

Monitored metric that determines if resource utilization is too low or high. If using EC2 metrics, consider enabling detailed monitoring for better scaling performance.

Target value

Instance warmup [Info](#)

seconds

☐ Disable scale in to create only a scale-out policy

Instance maintenance policy [Info](#)

Control your Auto Scaling group's availability during instance replacement events. This includes health checks, instance refreshes, maximum instance lifetime features and events that happen automatically to keep your group balanced, called rebalancing events.

Choose a replacement behavior depending on your availability requirements

Mixed behavior

☒ No policy

For rebalancing events, new instances will launch before terminating others. For all other events, instances terminate and launch at the same time.

Prioritize availability

☐ Launch before terminating

Launch new instances and wait for them to be ready before terminating others. This allows you to go above your desired capacity by a given percentage and may temporarily increase costs.

Control costs

☐ Terminate and launch

Terminate and launch instances at the same time. This allows you to go below your desired capacity by a given percentage and may temporarily reduce availability.

Flexible

☐ Custom behavior

Set custom values for the minimum and maximum amount of available capacity. This gives you greater flexibility in setting how far below and over your desired capacity EC2 Auto Scaling goes when replacing instances.

Additional capacity settings

Capacity Reservation preference [Info](#)

Select whether you want Auto Scaling to launch instances into an existing Capacity Reservation or Capacity Reservation resource group.

☒ Default

Auto Scaling uses the Capacity Reservation preference from your launch template.

☐ None

Instances will not be launched into a Capacity Reservation.

☐ Capacity Reservations only

Instances will only be launched into a Capacity Reservation. If capacity isn't available, the instances fail to launch.

☐ Capacity Reservations first

Instances will attempt to launch into a Capacity Reservation first. If capacity isn't available, instances will run in On-Demand capacity.

Additional settings

Instance scale-in protection

If protect from scale in is enabled, newly launched instances will be protected from scale in by default.

☐ Enable instance scale-in protection

Monitoring [Info](#)

Additional settings

Instance scale-in protection

If protect from scale in is enabled, newly launched instances will be protected from scale in by default.

☐ Enable instance scale-in protection

Monitoring [Info](#)

☐ Enable group metrics collection within CloudWatch

Default instance warmup [Info](#)

The amount of time that CloudWatch metrics for new instances do not contribute to the group's aggregated instance metrics, as their usage data is not reliable yet.

☐ Enable default instance warmup

Auto Scaling group deletion protection - *new* [Info](#)

Configure how this Auto Scaling group can be deleted.

☒ None (default)

The Auto Scaling group can be deleted.

☐ Prevent force deletion

The Auto Scaling group cannot be force deleted. To delete the group, all instances must be detached or terminated, and all warm pools must be deleted.

☐ Prevent all deletion

The Auto Scaling group cannot be deleted.

Cancel

Skip to review

Previous

Next

Review [Info](#)

Step 1: Choose launch template

[Edit](#)

Group details

Auto Scaling group name
Dynamic-CloudStorage-ASG

Launch template

Launch template
[Elastic-Storage-ASG-Template](#) [l2](#)
lt-08c8c1bb9146f7956

Version
Default

Description

Step 2: Choose instance launch options

[Edit](#)

Network

VPC
[vpc-02492443d8bcb7b0e](#) [l2](#)

Availability Zones and subnets

Availability Zone	Subnet	Subnet CIDR range
us-east-1-az4 (us-east-1c)	subnet-0b1657fa09eff134b l2	172.31.16.0/20
us-east-1-az1 (us-east-1a)	subnet-0f6dd550b930586bd l2	172.31.0.0/20
us-east-1-az2 (us-east-1b)	subnet-0866de87910fe34ae l2	172.31.80.0/20
us-east-1-az6 (us-east-1d)	subnet-0cbb951ae534cf4ec l2	172.31.32.0/20

Instance type requirements

This Auto Scaling group will adhere to the launch template.

Step 3: Integrate with other services

[Edit](#)

Load balancing

Load balancer
-

VPC Lattice integration options

VPC Lattice target groups
-

Application Recovery Controller (ARC) zonal shift

ARC zonal shift
Disabled

Health checks

Health check type
EC2, EBS

Health check grace period
300 seconds

Step 4: Configure group size and scaling policies

[Edit](#)

Group size

Desired capacity
1

Desired capacity type
Units (number of instances)

Scaling

Minimum desired capacity
1

Maximum desired capacity
4

Target tracking policy
Policy type
Target tracking scaling

Scaling policy name
Target Tracking Policy

Execute policy when
As required to maintain Average CPU utilization at 60

Take the action
Add or remove capacity units as required

Instances need
300 seconds to warm up before including in metric

Scale in
Enabled

Instance maintenance policy

Replacement behavior
No policy

Min healthy percentage
-

Max healthy percentage
-

Additional settings

Instance scale-in protection
Disabled

Monitoring
Disabled

Default instance warmup
Disabled

Auto Scaling group deletion protection
None (default)

Capacity Reservation preference

Preference

Default

Capacity Reservation IDs

-

Resource Groups

-

Step 5: Add notifications

Edit

Notifications

No notifications

Step 6: Add tags

Edit

Tags (0)

Key

Value

Tag new instances

No tags

Preview code

Cancel

Previous

Create Auto Scaling group

Create alarm

United States (IL, Virginia)

Swagata Sinha

Step 1

Specify metric and conditions

Step 2

Configure actions

Step 3

Add alarm details

Step 4

Preview and create

Specify metric and conditions

Edit

Metric

Graph

This alarm will trigger when the blue line goes above the red line for 1 datapoints within 1 minute.

Bytes

1,194

1000000

550k

0

16:00

16:30

17:00

17:30

18:00

18:30

NetworkIn

Namespace

AWS/EC2

Metric name

NetworkIn

AutoScalingGroupName

Dynamic-CloudStorage-ASG

Statistic

Average

Period

1 minute

Conditions

Threshold type

Static

Use a value as a threshold

Anomaly detection

Use a band as a threshold

Whenever NetworkIn is...

Define the alarm condition.

Greater

> threshold

Greater/Equal

>= threshold

Lower/Equal

<= threshold

Lower

< threshold

than...

Define the threshold value.

1000000

Must be a number.

Configure actions

Notification

Alarm state trigger

Define the alarm state that will trigger this action.

[Remove](#)

☒ **In alarm**

The metric or expression is outside of the defined threshold.

☐ **OK**

The metric or expression is within the defined threshold.

☐ **Insufficient data**

The alarm has just started or not enough data is available.

Send a notification to the following SNS topic

Define the SNS (Simple Notification Service) topic that will receive the notification.

☒ **Select an existing SNS topic**

☐ Create new topic

☐ Use topic ARN to notify other accounts

Send a notification to...

Q Dynamic_ASG

X

Only topics belonging to this account are listed here. All persons and applications subscribed to the selected topic will receive notifications.

Email (endpoints)

swagatasinha2108@gmail.com - [View in SNS Console](#)

[Add notification](#)

Lambda action

[Add Lambda action](#)

Auto Scaling action

[Add Auto Scaling action](#)

EC2 action

This action is only available for EC2 Per-Instance Metrics.

[Add EC2 action](#)

Systems Manager action [Learn more](#)

This action will create an Incident or Opsitem in Systems Manager when the alarm is **In alarm** state.

[Add Systems Manager action](#)

Investigation action - **new** [Info](#)


This action will create an Investigation when the alarm is **In alarm** state.

[Add investigation action](#)

Step 1: Specify metric and conditions

Metric

Graph
This alarm will trigger when the blue line goes above the red line for 1 datapoints within 5 minutes.



Namespace
AWS/EC2

Metric name
NetworkIn

AutoScalingGroupName
Dynamic-CloudStorage-ASG

Statistic
Average

Period
5 minutes

Conditions

Threshold type
Static

Whenever **NetworkIn** is
Greater (>)

than...
1000000

[Dynamic-CloudStorage-ASG](#) > Dynamic scaling policy

Edit dynamic scaling policy

Policy type
Step scaling

Scaling policy name
ASG-NetworkIn-ScaleOut

CloudWatch alarm
Choose an alarm that can scale capacity whenever:
ASG-NetworkIn-High-ScaleOut
[Create a CloudWatch alarm](#)
breaches the alarm threshold: NetworkIn > 1000000 for 1 consecutive periods of 300 seconds for the metric dimensions:
AutoScalingGroupName = Dynamic-CloudStorage-ASG

Take the action
Add
1 capacity units when 1000C <= NetworkIn < +infinity
[Add step](#)

Instance warmup | [Info](#)
300 seconds

Specify metric and conditions

Metric

Graph

This alarm will trigger when the blue line goes below the red line for 2 datapoints within 2 minutes.

Bytes

330k

300000

165k

0

17:0017:3018:0018:3019:0019:30

NetworkIn

Namespace

AWS/EC2

Metric name

NetworkIn

AutoScalingGroupName

Dynamic-CloudStorage-ASG

Statistic

Average

Period

1 minute

Edit

Conditions

Threshold type

Static

Use a value as a threshold

Anomaly detection

Use a band as a threshold

Whenever NetworkIn is...

Define the alarm condition.

Greater

> threshold

Greater/Equal

>= threshold

Lower/Equal

<= threshold

Lower

< threshold

than...

Define the threshold value.

300000

Must be a number.

Step 1: Specify metric and conditions

Metric

Graph

This alarm will trigger when the blue line goes below the red line for 2 datapoints within 2 minutes.

Bytes

330k

300000

165k

0

17:0017:3018:0018:3019:0019:30

NetworkIn

Namespace

AWS/EC2

Metric name

NetworkIn

AutoScalingGroupName

Dynamic-CloudStorage-ASG

Statistic

Average

Period

1 minute

Conditions

Threshold type

Static

Whenever NetworkIn is

Lower (<)

than...

300000

Create dynamic scaling policy

Policy type

Step scaling

Scaling policy name

ASG-NetworkIn-ScaleIn (Remove Instance)

CloudWatch alarm

Choose an alarm that can scale capacity whenever:

ASG-NetworkIn-Low-ScaleIn (Remove Instance)



[Create a CloudWatch alarm](#)

breaches the alarm threshold: NetworkIn < 300000 for 2 consecutive periods of 60 seconds for the metric dimensions:

AutoScalingGroupName = Dynamic-CloudStorage-ASG

Take the action

Remove

1

capacity units

when

30000

>=

NetworkIn

>

-infinity

Add step

Dynamic-CloudStorage-ASG



ASG-NetworkIn-ScaleIn (Remove Instance)

Policy type
Step scaling

Enabled or disabled
Enabled

Execute policy when
No alarm selected

Take the action
Remove 1 capacity units when +infinity <= Metric name < 0

ASG-NetworkIn-ScaleOut

Policy type
Step scaling

Enabled or disabled
Enabled

Execute policy when
ASG-NetworkIn-Low-ScaleIn
breaches the alarm threshold: NetworkIn > 1000000 for 2 consecutive periods of 60 seconds for the metric dimensions:
AutoScalingGroupName = Dynamic-CloudStorage-ASG

Take the action
Add 1 capacity units when 1000000 <= NetworkIn < +infinity

Instances need
300 seconds to warm up after each step

Target Tracking Policy

Policy type
Target tracking scaling

Enabled or disabled
Enabled

Execute policy when
As required to maintain Average CPU utilization at 60

Take the action
Add or remove capacity units as required

Instances need
300 seconds to warm up before including in metric

Elastic Compute Cloud	USD 0.00
No Region	(USD 6.78)
Asia Pacific (Mumbai)	USD 0.01
US East (N. Virginia)	USD 6.77
Amazon Elastic Compute Cloud running Linux/UNIX	USD 5.10
\$0.0104 per On Demand Linux t3.micro Instance Hour	1.129 Hrs USD 0.01
\$0.0116 per On Demand Linux t2.micro Instance Hour	438.587 Hrs USD 5.09
Amazon Elastic Compute Cloud running Windows	USD 0.01
\$0.0196 per On Demand Windows t3.micro Instance Hour	0.474 Hrs USD 0.01
Amazon Elastic Compute Cloud WindowsT3CPUCredits	USD 0.01
\$0.096 per vCPU-Hour of T3 CPU Credits	0.104 vCPU-Hours USD 0.01
EBS	USD 1.65

Billing and Cost Management

Bills

Filter by service name or region name

1

Billing and Cost Management

Billing View

New

Home

Getting Started

Dashboards

Billing and Payments

Bills

Payments

Credits

Purchase Orders

Cost and Usage Analysis

Cost Explorer

Cost Explorer Saved Reports

Description

Usage Quantity

Amount in USD

No Region		(USD 6.78)
Asia Pacific (Mumbai)		USD 0.01
US East (N. Virginia)		USD 6.77
Amazon Elastic Compute Cloud running Linux/UNIX		USD 5.10
\$0.0104 per On Demand Linux t3.micro Instance Hour	1.129 Hrs	USD 0.01
\$0.0116 per On Demand Linux t2.micro Instance Hour	438.587 Hrs	USD 5.09
Amazon Elastic Compute Cloud running Windows		USD 0.01
\$0.0196 per On Demand Windows t3.micro Instance Hour	0.474 Hrs	USD 0.01
Amazon Elastic Compute Cloud WindowsT3CPUCredits		USD 0.01
\$0.096 per vCPU-Hour of T3 CPU Credits	0.104 vCPU-Hours	USD 0.01

Filter activity history

Filter by a date and time range

1 2 >

Status	Description	Cause	Start time	End time
Successful	Terminating EC2 instance: i-01380280c77f47c89	At 2026-02-19T18:52:25Z a monitor alarm ASG-NetworkIn-Low-ScaleIn (Remove Instance) in state ALARM triggered policy ASG-NetworkIn-ScaleIn (Remove Instance) changing the desired capacity from 2 to 1. At 2026-02-19T18:52:27Z an instance was taken out of service in response to a difference between desired and actual capacity, shrinking the capacity from 2 to 1. At 2026-02-19T18:52:28Z instance i-01380280c77f47c89 was selected for termination.	2026 February 20, 12:22:28 AM +05:30	2026 February 20, 12:23:10 AM +05:30
Successful	Launching a new EC2 instance: i-01380280c77f47c89	At 2026-02-19T18:46:49Z a monitor alarm ASG-NetworkIn-High-ScaleOut in state ALARM triggered policy ASG-NetworkIn-ScaleOut changing the desired capacity from 1 to 2. At 2026-02-19T18:47:04Z an instance was started in response to a difference between desired and actual capacity, increasing the capacity from 1 to 2.	2026 February 20, 12:17:06 AM +05:30	2026 February 20, 12:22:10 AM +05:30
Successful	Launching a new EC2 instance: i-084dfe8ee5f848c66	At 2026-02-15T06:15:00Z an instance was launched in response to an unhealthy instance needing to be replaced.	2026 February 15, 11:45:02 AM +05:30	2026 February 15, 11:45:07 AM +05:30
Successful	Terminating EC2 instance: i-0d5b18d8da4c2ddb4	At 2026-02-15T06:15:00Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.	2026 February 15, 11:45:00 AM +05:30	2026 February 15, 11:45:02 AM +05:30
Successful	Launching a new EC2 instance: i-0d5b18d8da4c2ddb4	At 2026-02-15T05:32:42Z an instance was launched in response to an unhealthy instance needing to be replaced.	2026 February 15, 11:02:44 AM +05:30	2026 February 15, 11:02:49 AM +05:30
Successful	Terminating EC2 instance: i-07beb88fb8b47a4a5	At 2026-02-15T05:32:42Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.	2026 February 15, 11:02:42 AM +05:30	2026 February 15, 11:03:24 AM +05:30
Successful	Launching a new EC2 instance: i-07beb88fb8b47a4a5	At 2026-02-15T05:02:38Z an instance was launched in response to an unhealthy instance needing to be replaced.	2026 February 15, 10:32:39 AM +05:30	2026 February 15, 10:32:44 AM +05:30
Successful	Terminating EC2 instance: i-00794d35f264352fe	At 2026-02-15T05:02:37Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.	2026 February 15, 10:32:37 AM +05:30	2026 February 15, 10:32:39 AM +05:30
Successful	Launching a new EC2 instance: i-00794d35f264352fe	At 2026-02-15T03:58:06Z an instance was launched in response to an unhealthy instance needing to be replaced.	2026 February 15, 09:28:08 AM +05:30	2026 February 15, 09:28:13 AM +05:30
Successful	Terminating EC2 instance: i-07df0d680be16ae73	At 2026-02-15T03:58:06Z an instance was taken out of service in response to an EC2 health check indicating it has been terminated or stopped.	2026 February 15, 09:28:06 AM +05:30	2026 February 15, 09:28:28 AM +05:30